



UNIVERSITE DE NOUAKCHOTT
FACULTE DES SCIENCES ET TECHNIQUES
FEPARTEMENT INFORMATIQUE

SDD



Machine Learning

Simple Linear Regression

Dr. EL BENANY Mohamed Mahmoud

30.11.24



<https://forms.office.com/r/yYL28qWb3m>

SUPERVISED MACHINE LEARNING

Régression

Simple linear regression

Multiple linear regression

Ridge and Lasso Regression

Polynomial regression

Decision Tree and Random Forest Regression

Classification

Logistic Regression Classification

Decision Tree Random Forest Classification VS Random

Naive Bayesian Classifiers for Ranking

Classification of support vector machines

Vector Machine Kernel Support

K-Nearest Neighbor Classifier

Evaluation of classification models



UNSUPERVISED LEARNING

Clustering

Dimensionality Reduction

Data Analysis

Artificial intelligence

Machine Learning

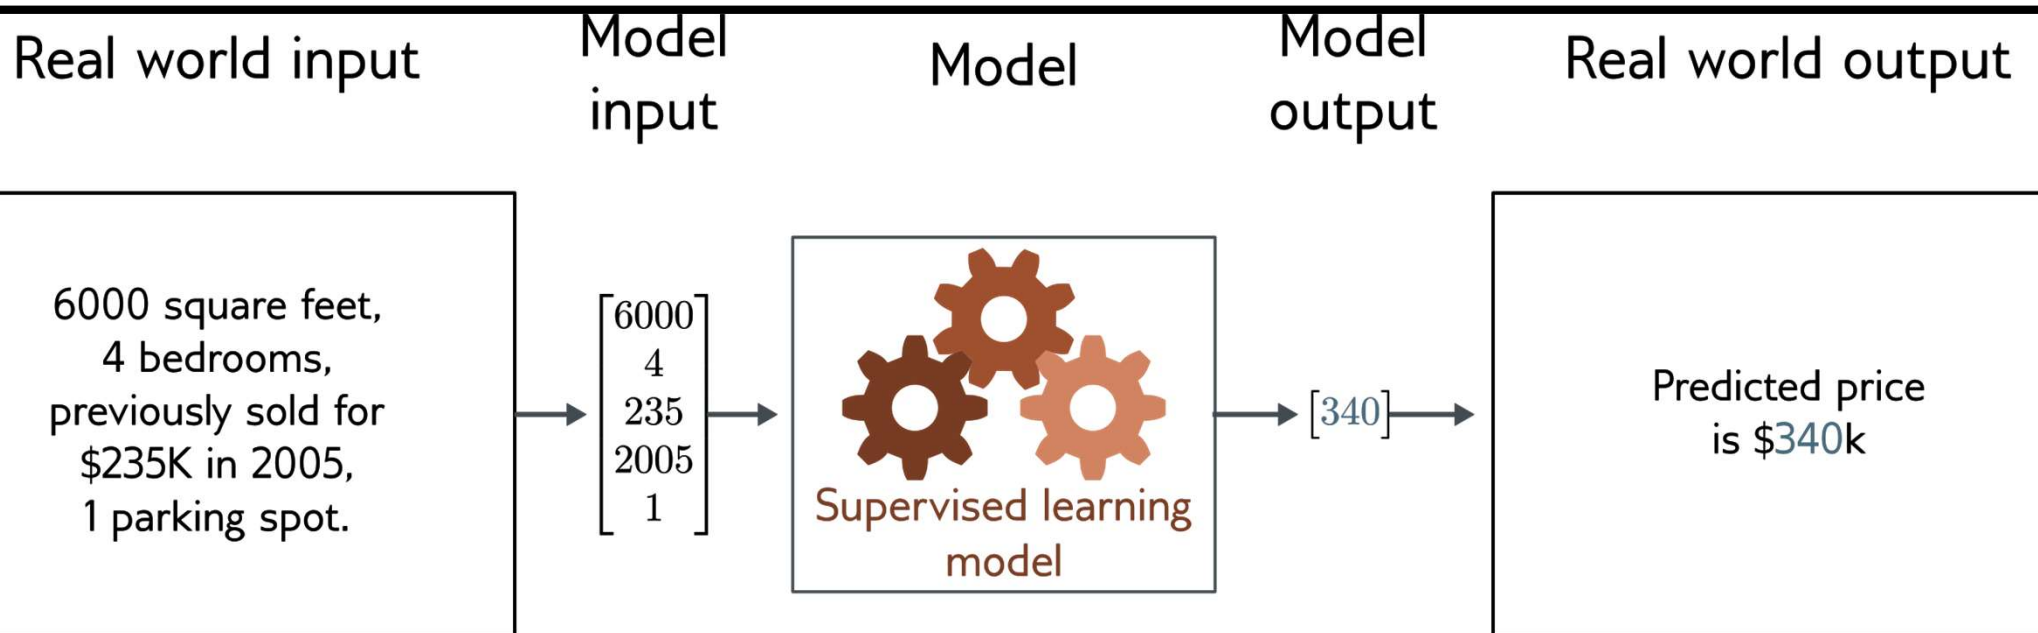
Supervised
learning

Unsupervised
Learning

Learning
By reinforcement

Deep learning

Regression



Univariate regression problem (one output, real value)

Supervised learning overview

Supervised learning model = mapping from one or more inputs to one or more outputs

- Model is a mathematical equation
- Computing the outputs from the inputs = **inference**
- Example:
 - Input is age and milage of secondhand Toyota Prius
 - Output is estimated price of car

Supervised learning overview

- **Supervised learning model** = mapping from one or more inputs to one or more outputs
- Model is a mathematical equation
- Computing the inputs from the outputs = **inference**
- Model also includes **parameters**
- Parameters affect outcome of equation

Supervised learning overview

- **Supervised learning model** = mapping from one or more inputs to one or more outputs
- ~~Model is a mathematical equation~~
- Model is a family of equations
- Computing the inputs from the outputs = **inference**
- Model also includes **parameters**
- Parameters affect outcome of equation

Supervised learning overview

- **Supervised learning model** = mapping from one or more inputs to one or more outputs
- Model is a family of equations
- Computing the inputs from the outputs = **inference**
- Model also includes **parameters**
- Parameters affect outcome of equation

Training a model = finding parameters that predict outputs “well” from inputs for a **training dataset** of input/output pairs

Supervised learning

- Overview
- **Notation**
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Notation

- Input:

x

Variables always Roman letters

- Output:

y

Normal = scalar
Bold = vector
Capital Bold = matrix

- Model:

$$y = \mathbf{f}[\mathbf{x}]$$

Functions always square brackets

Normal = returns scalar
Bold = returns vector
Capital Bold = returns matrix

Supervised learning

- Overview
- **Notation**
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Notation example

- Input:

$$\mathbf{x} = \begin{bmatrix} \text{age} \\ \text{mileage} \end{bmatrix}$$

← Structured or tabular data

- Output:

$$y = [\text{price}]$$

- Model:

$$y = f[\mathbf{x}]$$

Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Model

- Parameters:

ϕ

Parameters always
Greek letters

- Model :

$$y = f[x, \phi]$$

Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Loss function

- Training dataset of I pairs of input/output examples:

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I$$

- Loss function or cost function measures how bad model is:

$$L[\phi, f[\mathbf{x}_i, \phi], \{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I]$$

or for short:

$$L[\phi]$$

← Returns a scalar that is smaller when model maps inputs to outputs better

Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Training

- Loss function:

$$L[\phi]$$

Returns a scalar that is smaller
when model maps inputs to
outputs better

- Find the parameters that minimize the loss:

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} [L[\phi]]$$

Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Testing

- To test the model, run on a separate **test dataset** of input / output pairs
- See how well it **generalizes** to new data

Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

1D Linear regression example

- Model:

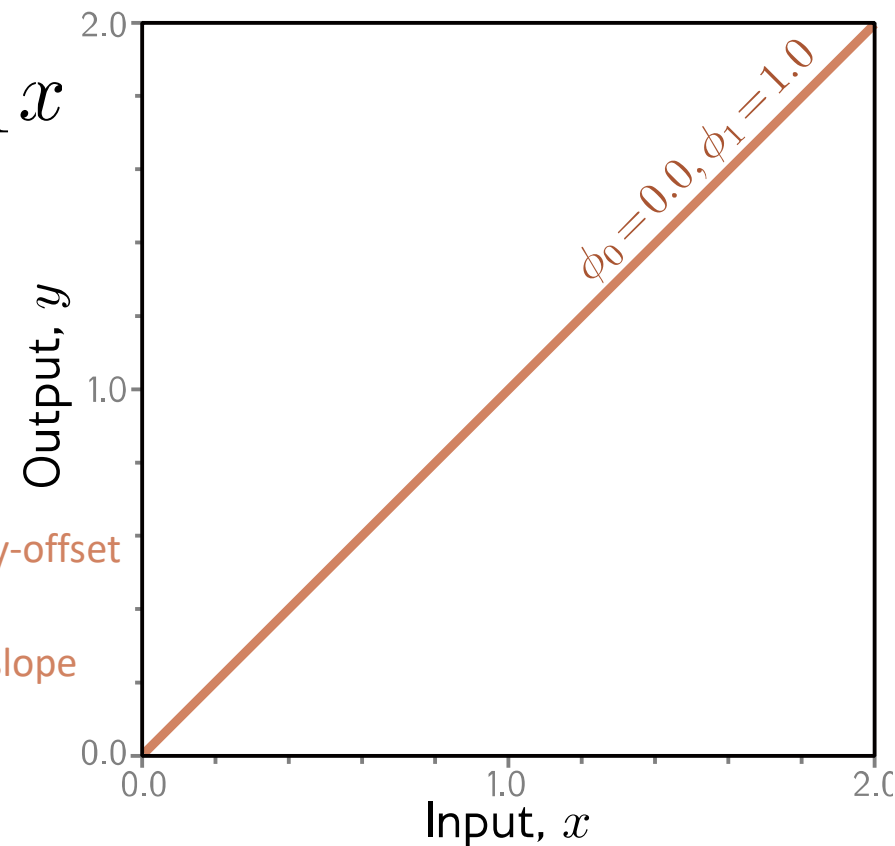
$$y = f[x, \phi]$$
$$= \phi_0 + \phi_1 x$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset

← slope



Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

1D Linear regression example

- Model:


$$y = f[x, \phi]$$
$$= \phi_0 + \phi_1 x$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset

← slope



Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

1D Linear regression example

- Model:


$$y = f[x, \phi] \\ = \phi_0 + \phi_1 x$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset


← slope



Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

1D Linear regression example



Supervised learning

- 1D Linear regression example


- Model
- Loss function
- Training
- Testing

Loss function:

$$\begin{aligned} L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”


1D Linear regression example



Supervised learning

• 1D Linear regression example

- Model
- Loss function
- Training
- Testing




1D Linear regression example

Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”



Supervised learning

- 1D Linear regression example

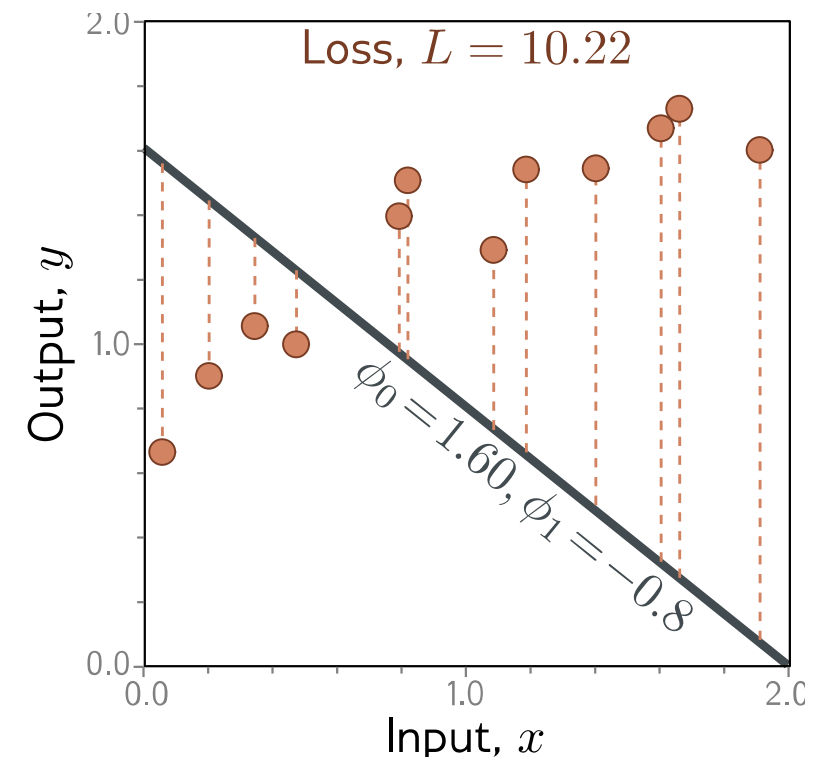

- Model
- Loss function
- Training
- Testing

1D Linear regression example

Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”



Supervised learning

- 1D Linear regression example



- Model
- Loss function
- Training
- Testing

1D Linear regression example

Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”



Supervised learning

• 1D Linear regression example

- Model
- Loss function
- Training
- Testing

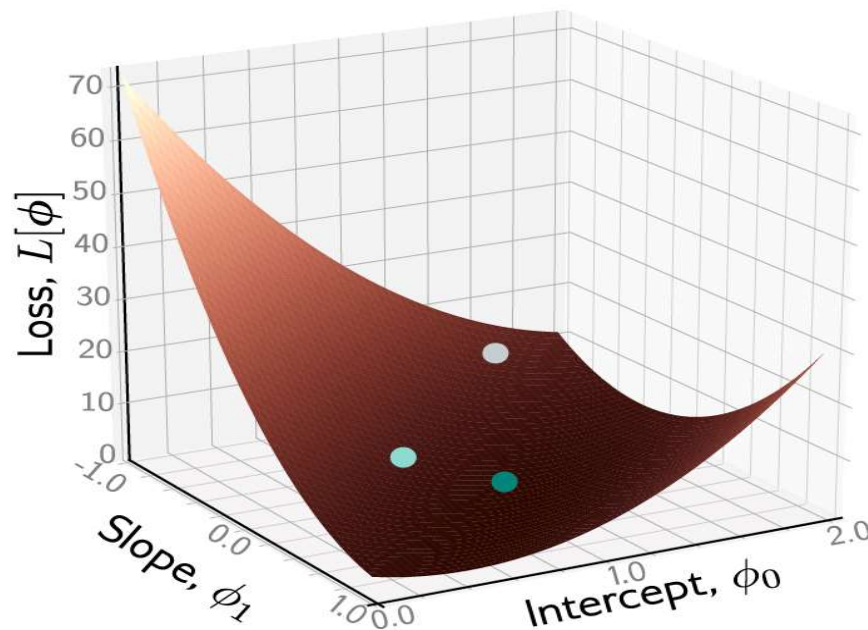
1D Linear regression example

Loss function:


$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”

a)



b)





Supervised learning

Example: 1D Linear regression training

- 1D Linear regression example

- Model
- Loss function
- Training
- Testing





Supervised learning

Example: 1D Linear regression training

- 1D Linear regression example

- Model
- Loss function
- Training
- Testing





Supervised learning

• 1D Linear regression example

- Model
- Loss function
- Training
- Testing

Example: 1D Linear regression training





Supervised learning

• 1D Linear regression example

- Model
- Loss function
- Training
- Testing

Example: 1D Linear regression training





Supervised learning

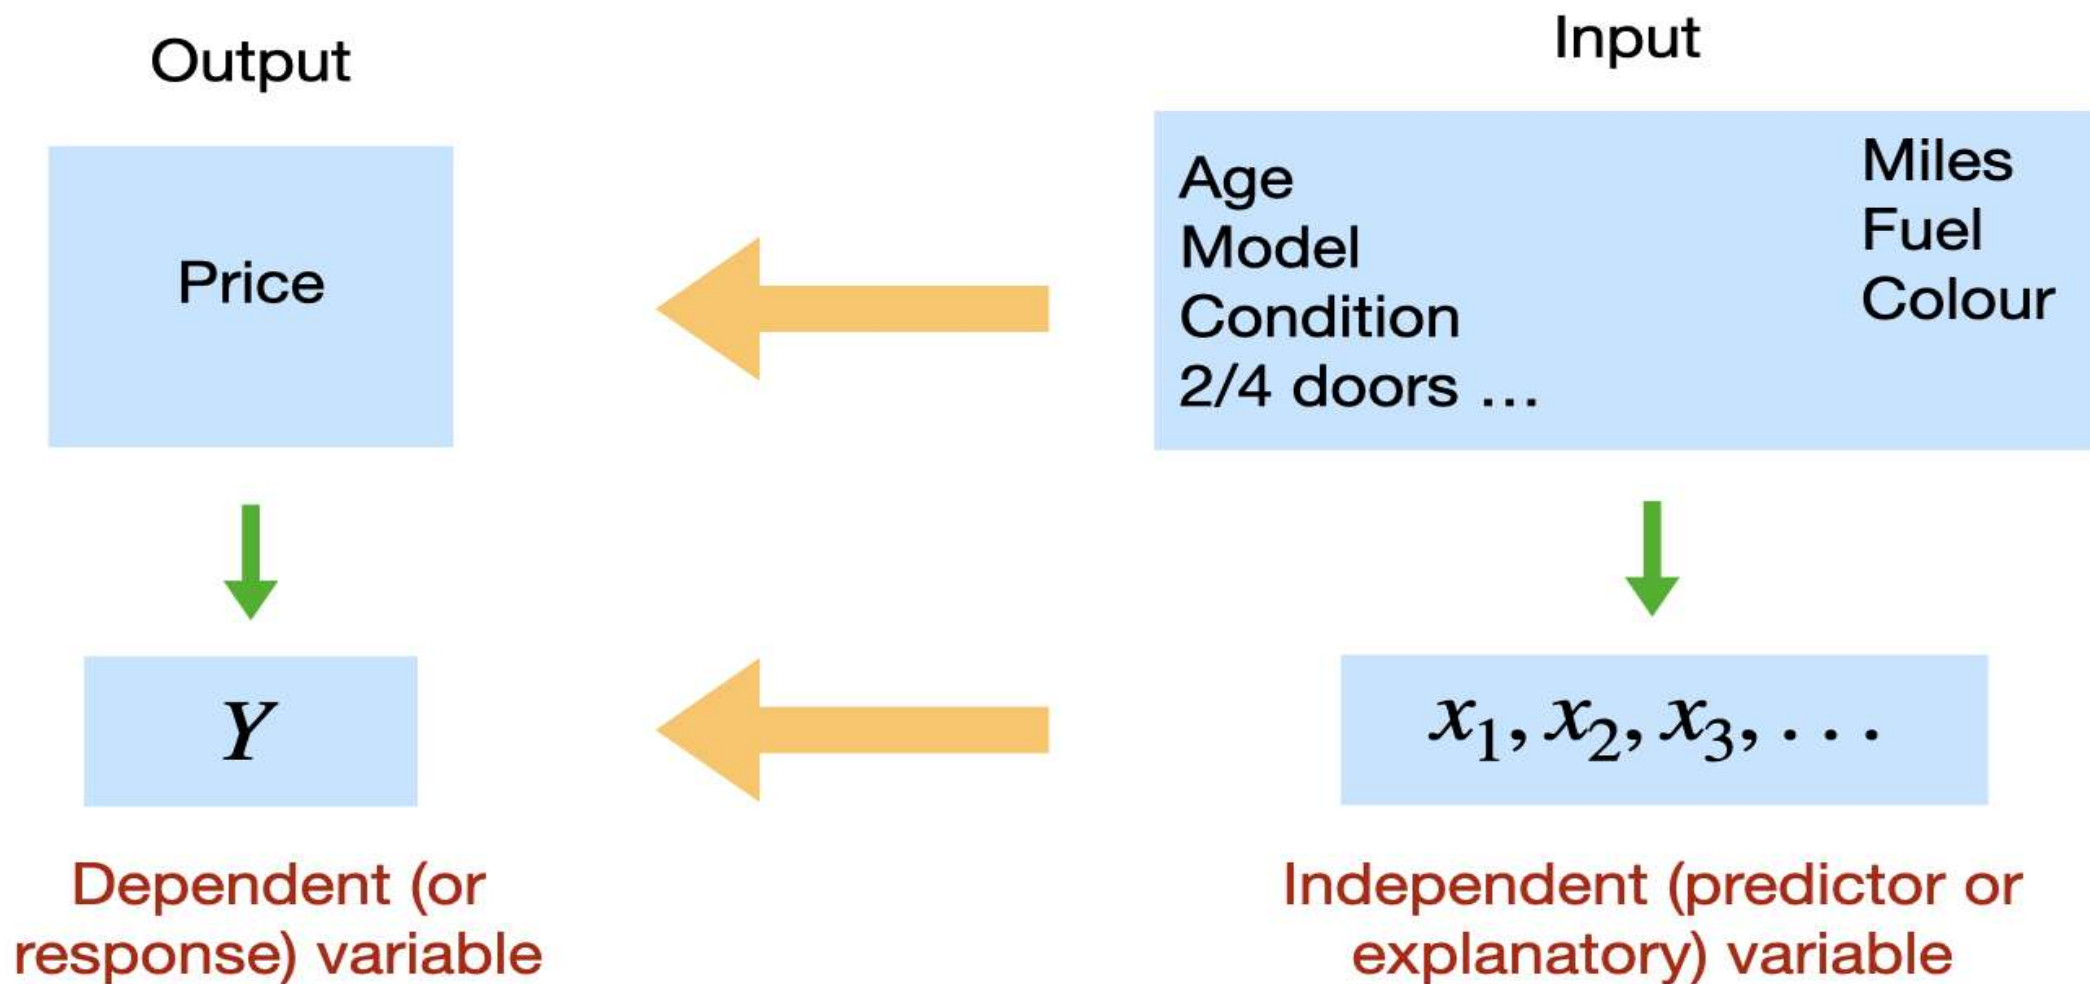
• 1D Linear regression example

- Model
- Loss function
- Training
- Testing

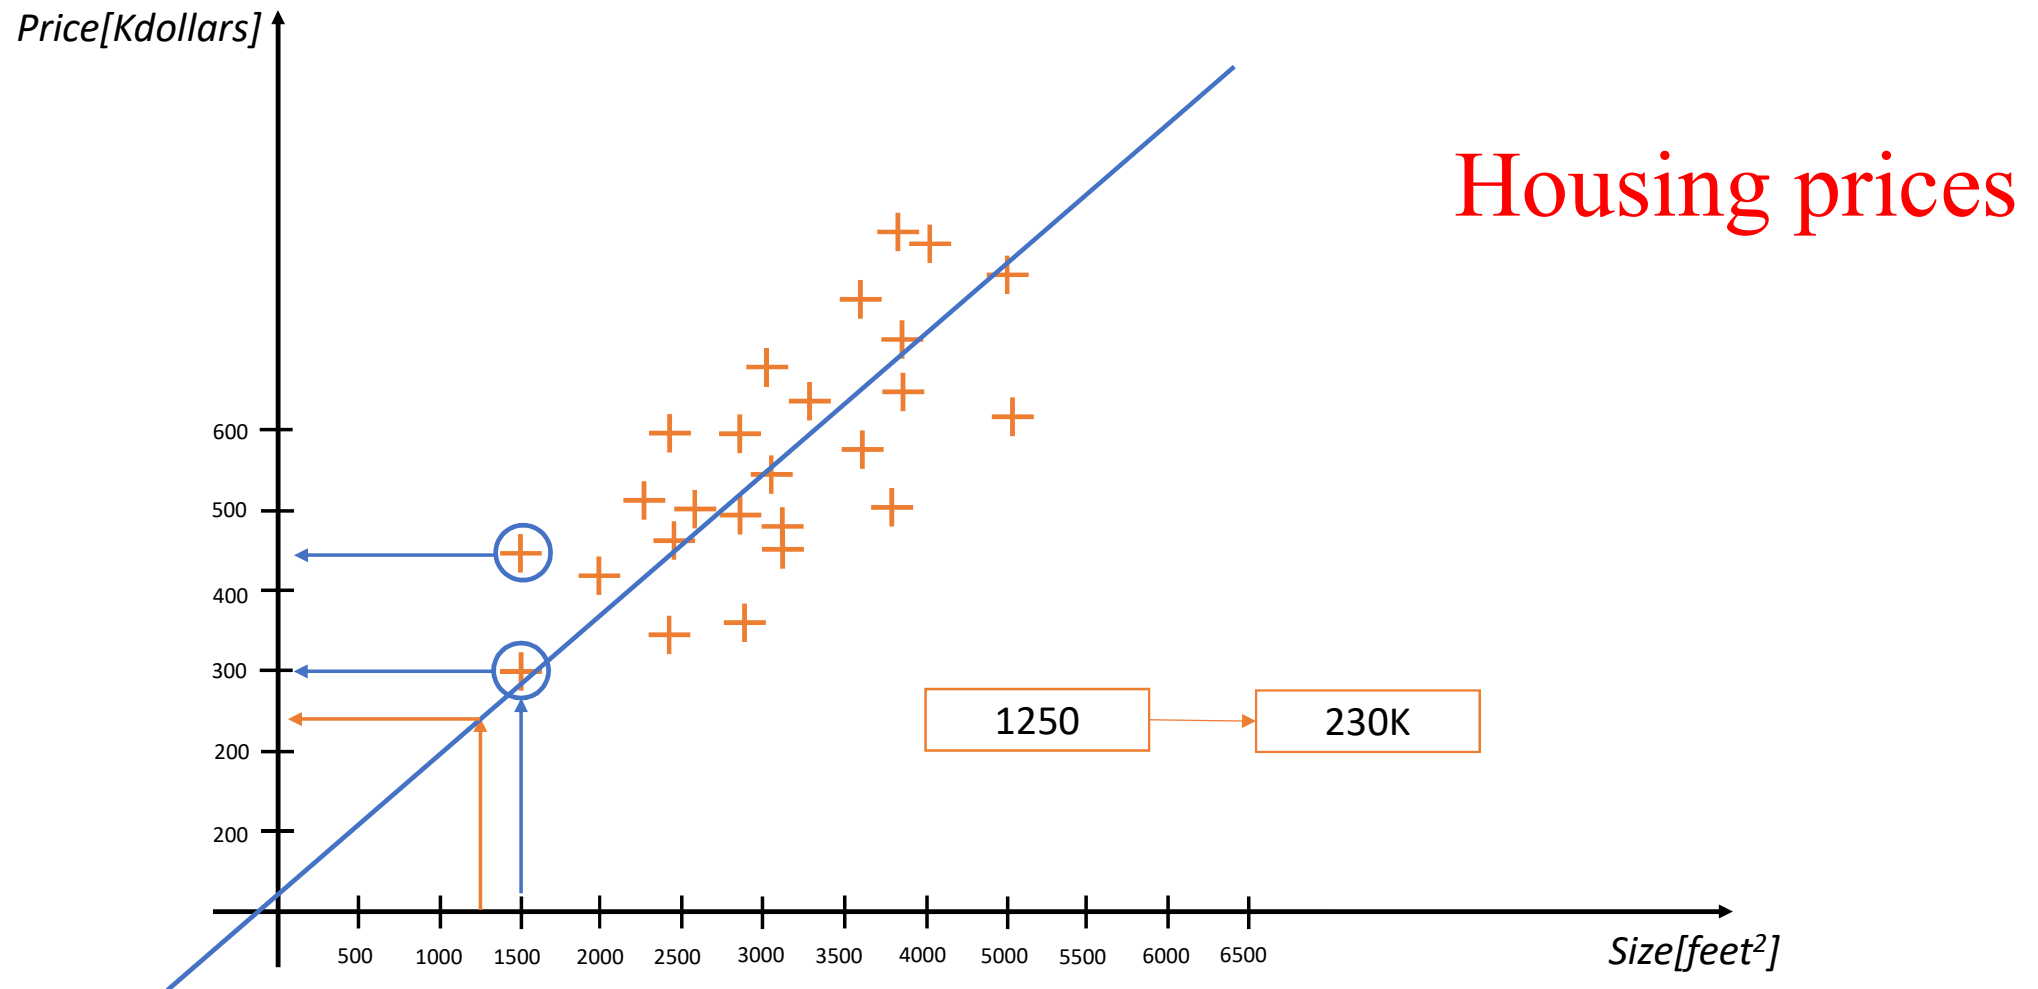
Example: 1D Linear regression training



Linear regression variables



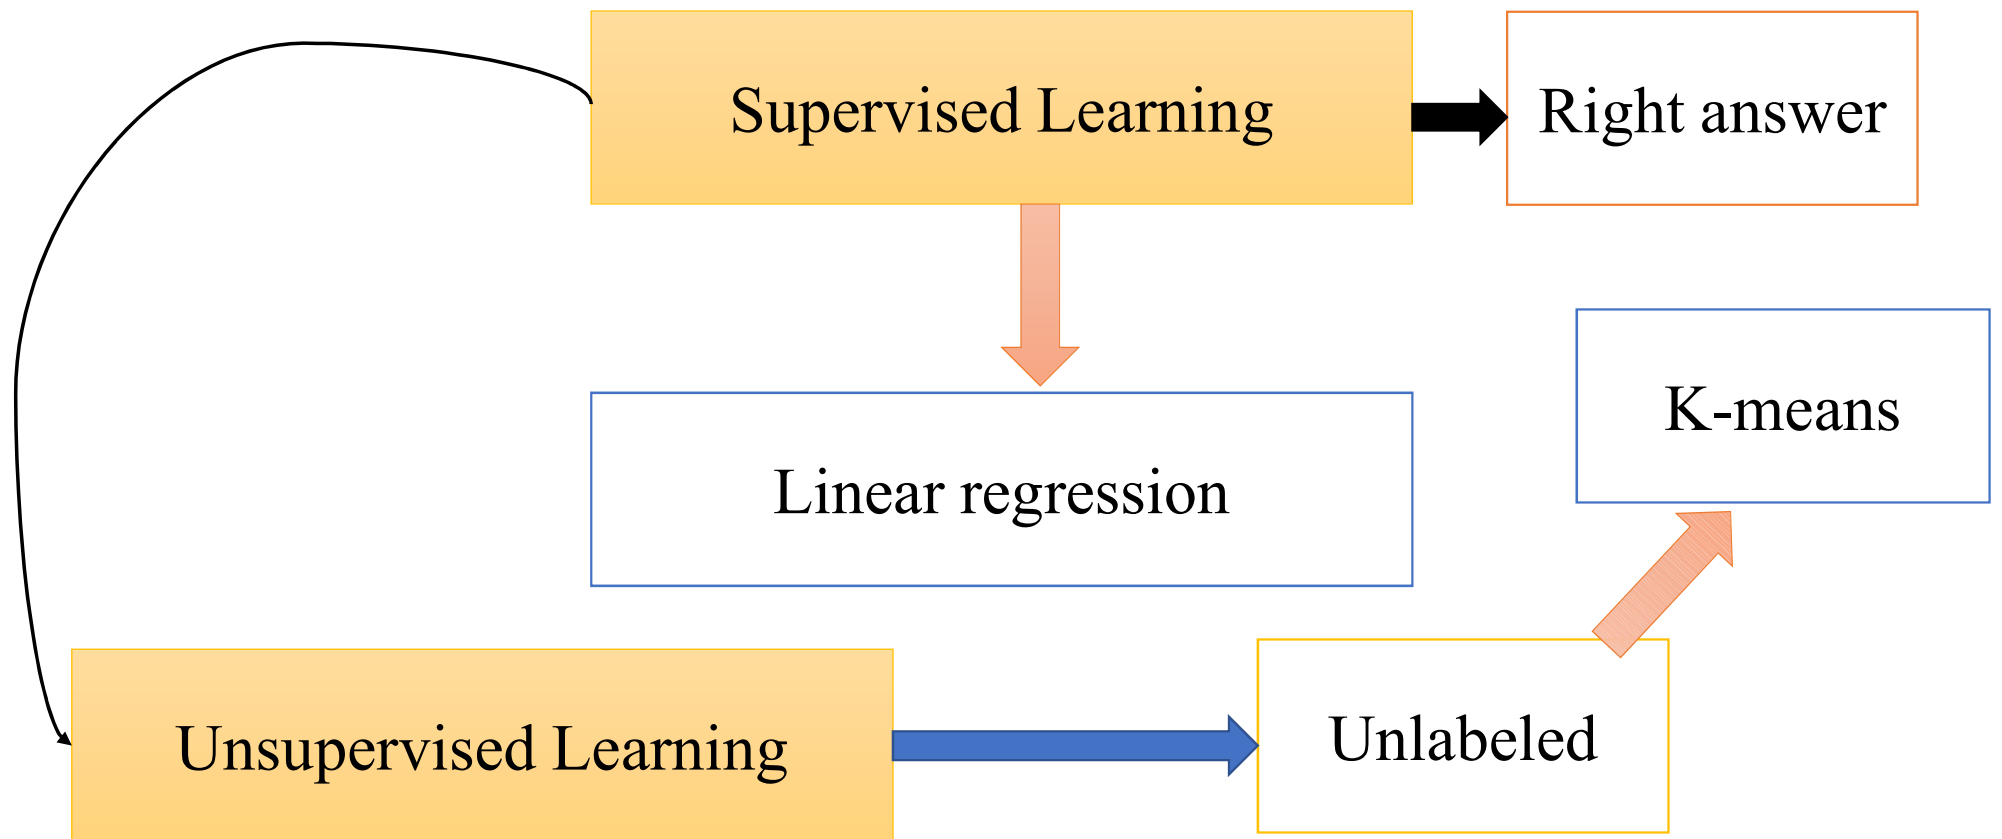
Simple Linear Regression



Supervised Learning

- Linear regression:
 - Real values: 1.2, 2.5, 3.8, 4.2
- Classification:
 - Discrete values: 1, 2, 3, 4

Machine Learning



Terminologies: Training Data

X

Y

➡ Input variable

➡ Output variable

➡ Features

➡ Target

➡ Independent variable

➡ Dependent variable

Size[feet ²]	<i>Price in K\$</i>
51000	600
48000	300
38000	403
11000	210
...	...

$X^i \& Y^i$

$X^1 = 51000$

$X^2 = 48000$

$Y^1 = 600 \dots$

Simple Linear Regression (Housing Data Analysis)

Training Data



Learning Algorithm



Size

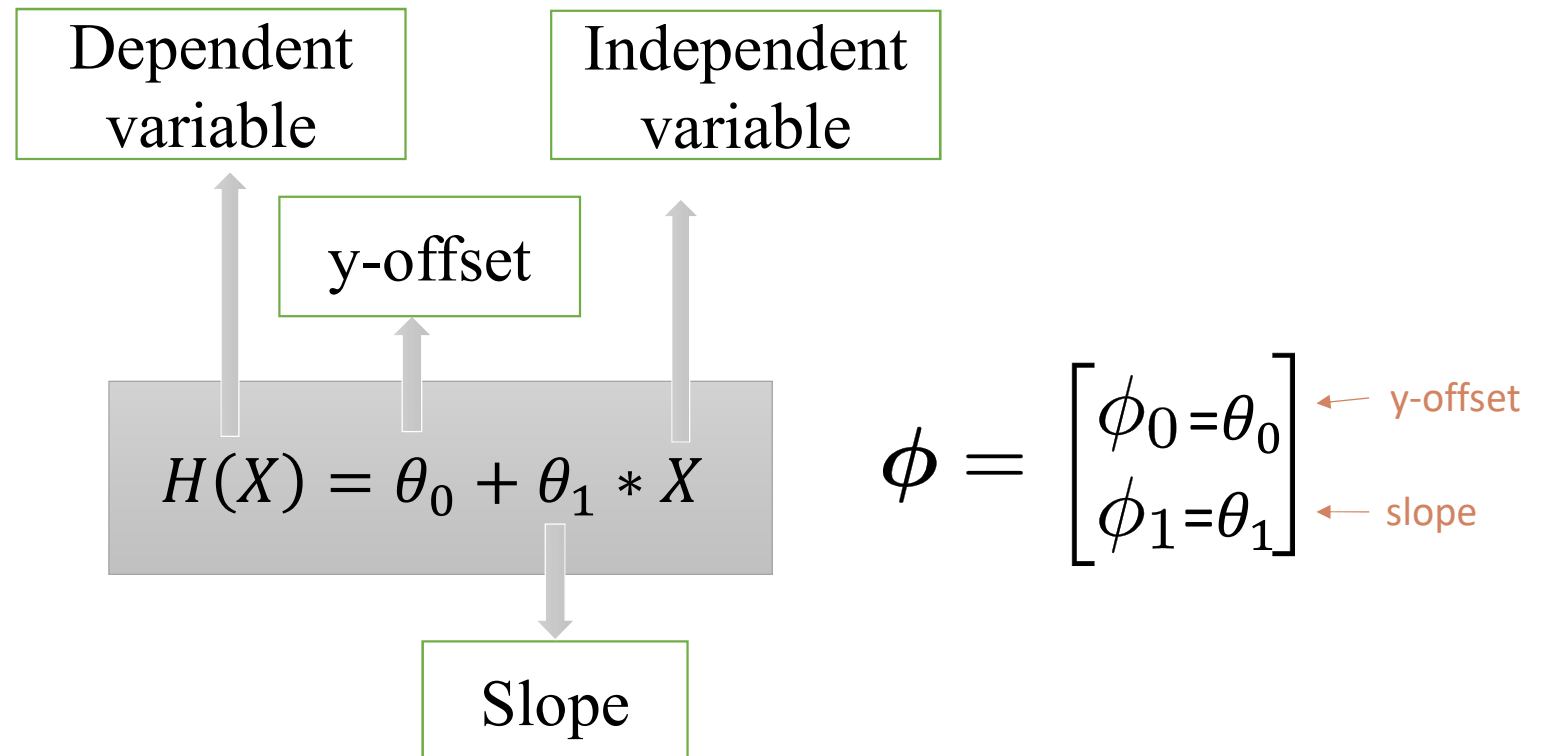
Function

H




Estimated Price


Univariate Linear Regression



Hypothesis Function (1)



Hypothesis Function (1)




$$H(X) = \theta_0 + \theta_1 * X$$

$$H(X) = 2 + 0$$

$$\theta_0 = 2$$

$$\theta_1 = 0$$

Hypothesis Function (2)




$$H(X) = 0 + 2 * X$$


$$\theta_0 = 0$$

$$\theta_1 = 2$$


Hypothesis Function (3)



Errors



All Errors...




ERROR

$$Error_1 = Y_1 - (\theta_0 + \theta_1 * X_1) \rightarrow Y_1 - H(X_1)$$

$$Error_2 = Y_2 - (\theta_0 + \theta_1 * X_2) \rightarrow Y_2 - H(X_2)$$

All Errors...



MSE

$$\frac{1}{2m} \sum_{i=1}^m [H(X^i) - Y^i]^2$$

Absolute Error Function?

Square Error Function

Cost function ⁴⁴


Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [H(X^i) - Y^i]^2$$

Exemple($\theta_1 = 1$)



$$H(X) = \theta_1 * X \quad \theta_1 = 1$$

$$H(X) = X$$





$$J(1) = \frac{1}{2 * 4} (0^2 + 0^2 + 0^2 + 0^2) = 0$$

Exemple($\theta_1 = 0.4$)




$$J(0.4) = \frac{1}{2 * 4} ((0.4 * 2 - 2)^2 + (0.4 * 3 - 3)^2 + \dots) = 2.43$$

Exemple($\theta_1 = 0$)



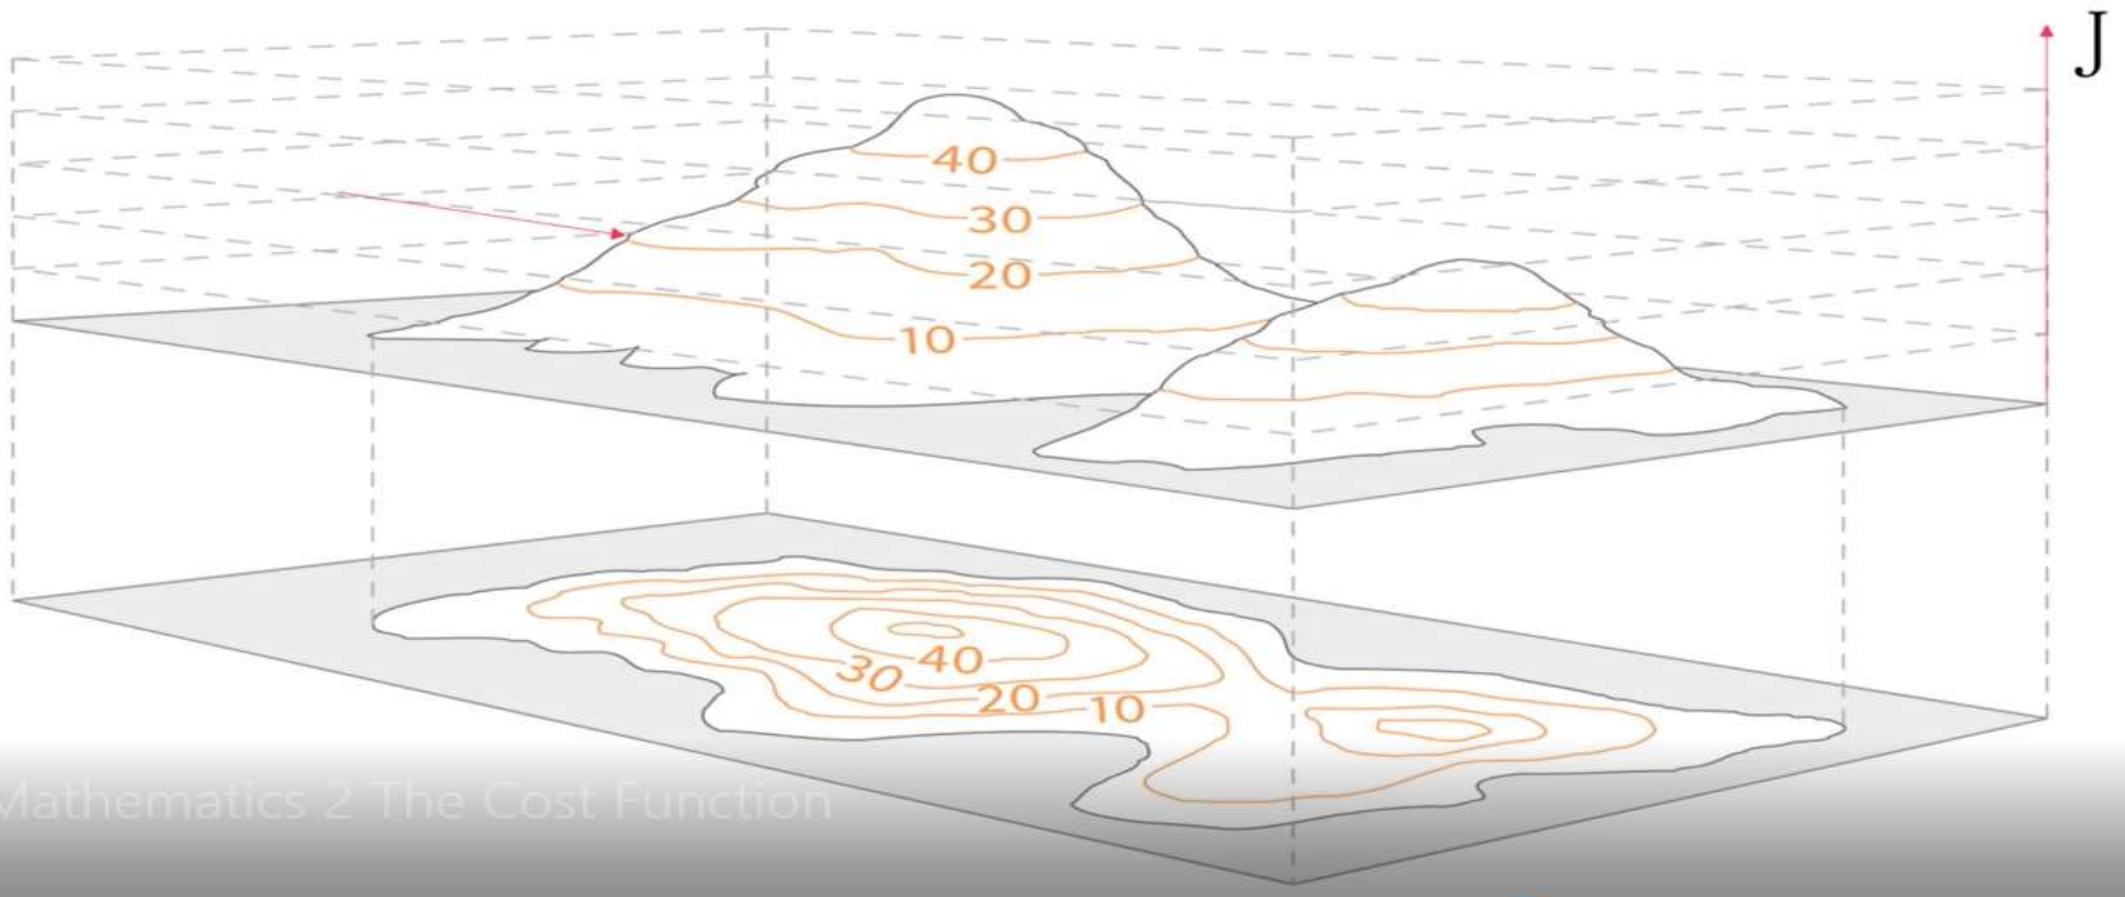
$$J(0) = \frac{1}{2 * 4} ((2)^2 + (3)^2 + (4)^2 + (5)^2) = 9$$

Cost function




$$H(X) = \theta_0 + \theta_1 * X$$



Cost function



Cost function





Cost function





$$H(X) = \theta_0 + \theta_1 * X$$

Cost function



$$H(X) = \theta_0 + \theta_1 * X$$

Cost function



$$H(X) = \theta_0 + \theta_1 * X$$

Optimization

Gradient Descent

- The Gradient: Vector indicates the direction in which the function is increasing by largest amount.

Multivariable Functions

$$F(X, Y)$$

$$\nabla \mathcal{F} = \frac{\partial \mathcal{F}}{\partial x} \hat{i} + \frac{\partial \mathcal{F}}{\partial y} \hat{j}$$


Gradient Descent

- The Gradient: Vector indicates the direction in which the function is increasing by largest amount.


Multivariable Functions

$$F(X, Y)$$

$$\nabla F = \frac{\partial F}{\partial x} \mathbf{i} + \frac{\partial F}{\partial y} \mathbf{j}$$



How GD Works




$$\theta \stackrel{=}{:=} \theta - \nabla J(\theta_0 + \theta_1)$$

Assignment $\stackrel{=}{=}$


Equality $\stackrel{==}{=}$

How GD Works



$$\theta := \theta - \nabla J(\theta_0 + \theta_1)$$

$$\theta_0 = 0$$



$$\theta_1 := \theta_1 - \nabla J(\theta_1)$$

How GD Works

$$\theta_1 := \theta_1 - \nabla J(\theta_1)$$


$$\theta_1 := 1.8$$

$$\theta_1 := 1.8 - \nabla J(\theta_1)$$


$$\theta_1 := 1.8 - \frac{dJ}{d\theta_1} \quad 0.4$$

$$\theta_1 := 1.4 - \frac{dJ}{d\theta_1} \quad 0.2$$

$$\theta_1 := 1.2 - \frac{dJ}{d\theta_1} \quad 0.1$$



Query 1 : What about the Initialization?



Query 2 How to Adjust the Speed of Algorithm?

$$\theta_1 := \theta_1 - \nabla J(\theta_1)$$


$$\theta_1 := 1.8$$

$$\theta_1 := 1.8 - \nabla J(\theta_1)$$

$$\theta_1 := 1.8 - 2 * \frac{dJ}{d\theta_1}$$

$$\theta_1 := 1.8 - 0.8$$

$$\theta_1 = 1$$



Query 2 How to Adjust the Speed of Algorithm?

$$\theta_1 \coloneqq \theta_1 - \alpha * \nabla J(\theta_1)$$

Query 2 How to Adjust the Speed of Algorithm?

$$\theta_1 := \theta_1 - \alpha \nabla J(\theta_1)$$

$$\theta_1 := 1.8$$

$$\theta_1 := 1.8 - \nabla J(\theta_1)$$


$$\theta_1 := 1.8 - 3 * \frac{dJ}{d\theta_1}$$

$$\theta_1 = 0.6$$

$$\theta_1 = 1.4$$

$$\theta_1 = 0.88$$

$$\theta_1 = 1.45$$




$\theta_1 = ? ?$



Query 2 How to Adjust the Speed of Algorithm


$$\theta_1 := \theta_1 - \alpha \nabla J(\theta_1)$$

Farther




larger


Closer





Smaller





Query 3 What if it Was a Non-Convex Function?



Global vs Local Optimum



Global vs Local Optimum



Polymerization Between Gradient and Hypothesis

Linear Hypothesis

$$H(X) = \theta_0 + \theta_1 * X$$

$$\frac{1}{2m} \sum_{i=1}^m [H(X^i) - Y^i]^2$$

Gradient Descent

$$\theta_1 := \theta_1 - \alpha \nabla J(\theta_1)$$

Repeat until convergence: {
 $\theta_i := \theta_i - \alpha \nabla J(\theta_0, \theta_1)$
[$i = 0, i = 1$]} }

Repeat until convergence: {
 $\theta_i := \theta_i - \alpha \frac{d}{d\theta_i} J(\theta_0, \theta_1)$
[$i = 0, i = 1$]} }

Polymerization Between Gradient and Hypothesis

$$\left\{ \begin{array}{l} \theta_i := \theta_i - \alpha \frac{d}{d\theta_i} J(\theta_0, \theta_1) \\ [i = 0, i = 1] \end{array} \right\}$$

$$\frac{d}{d\theta_0} J(\theta_0, \theta_1) = \frac{d}{d\theta_0} \left(\frac{1}{2m} \sum_{i=1}^m [H(X^i) - Y^i]^2 \right)$$

$$= \frac{d}{d\theta_0} \left(\frac{1}{2m} \sum_{i=1}^m [\theta_0 + \theta_1 * X - Y^i]^2 \right)$$

$$= \left(\frac{1}{m} \sum_{i=1}^m [\theta_0 + \theta_1 * X - Y^i] \cdot 1 \right) = \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i]$$

$i = 0$

Polymerization Between Gradient and Hypothesis

$$\left\{ \begin{aligned} \theta_i &:= \theta_i - \alpha \frac{d}{d\theta_i} J(\theta_0, \theta_1) \\ [i = 0, i = 1] \end{aligned} \right\}$$

$$\frac{d}{d\theta_1} J(\theta_0, \theta_1) = \frac{d}{d\theta_1} \left(\frac{1}{2m} \sum_{i=1}^m [H(X^i) - Y^i]^2 \right)$$

$$= \frac{d}{d\theta_1} \left(\frac{1}{2m} \sum_{i=1}^m [\theta_0 + \theta_1 * X - Y^i]^2 \right)$$

$$= \left(\frac{1}{m} \sum_{i=1}^m [\theta_0 + \theta_1 * X - Y^i] \cdot X \right) = \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i] \cdot X$$

$i = 1$

Polymerization Between Gradient and Hypothesis

Repeat until convergence: {
{

$$\theta_0 := \theta_0 - \alpha \frac{d}{d\theta_0} J(\theta_0, \theta_1)$$

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_0, \theta_1)$$

}

Repeat until convergence: {
{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i]$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i] \cdot X$$

}

Polymerization Between Gradient and Hypothesis

Repeat until convergence: {

$$\theta_0 := \theta_0 - \alpha \frac{d}{d\theta_0} J(\theta_0, \theta_1)$$

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_0, \theta_1)$$

}

Together

θ_0

θ_1

50

0.5

60

0.5

Repeat until convergence: {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i]$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i] \cdot X$$

}

Repeat until convergence: {

{

$$T_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i]$$

$$T_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [H(X^i) - Y^i] \cdot X$$

$$\theta_0 := T_0$$

$$\theta_1 := T_1$$

}

TP: Required Libraries and Importing Data

1. Pandas
2. Numpy
3. Matplotlib

Error:

- Unicode Error(EBCDIC)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

iloc Function

```
# Chargement des données depuis un fichier CSV et séparation en variables indépendantes (indepX) et dépendantes (depY)
dataset = pd.read_csv("Datasets/Housing_Data.csv")
indepX = dataset.iloc[:, 0] # Variable indépendante (ex : surface de la maison)
depY = dataset.iloc[:, 1]  # Variable dépendante (ex : prix de la maison)
```

Handling Data

Splitting Data into Train and Test Sets

```
# Division des données en ensembles d'entraînement et de test (80% entraînement, 20% test)
indepX_train, indepX_test, depY_train, depY_test = train_test_split(indepX, depY, test_size=0.2, random_state=42)
```

GD Function

```
# Fonction de descente de gradient pour optimiser les paramètres du modèle linéaire
def gradientDescent(indepX, depY, init_theta, learning_rate, num_iterations):
    # Mise à jour itérative des paramètres theta pour minimiser l'erreur
    theta = init_theta
    for i in range(num_iterations):
        theta = grad(indepX, depY, theta, learning_rate)
    return theta
```

Gradient Function

```
# Fonction pour calculer le gradient des paramètres à chaque étape
def grad(indepX, depY, curr_theta, learning_rate):
    # Calcul du gradient basé sur la dérivée partielle de la fonction de coût
    grad = np.zeros(2) # Gradient pour theta[0] (biais) et theta[1] (poids)
    new_theta = curr_theta
    m = len(indepX)    # Nombre de données
    for i in range(m):
        x = indepX.iloc[i] # Correction : Utiliser `.iloc` pour éviter les erreurs avec Series
        y = depY.iloc[i]   # Correction : Idem
        grad[0] += (-1/m) * (y - (curr_theta[0] + curr_theta[1] * x))
        grad[1] += (-1/m) * x * (y - (curr_theta[0] + curr_theta[1] * x))
    # Mise à jour des paramètres en fonction du gradient
    new_theta = np.zeros(2)
    temp0 = curr_theta[0] - learning_rate * grad[0]
    temp1 = curr_theta[1] - learning_rate * grad[1]
    new_theta[0] = temp0
    new_theta[1] = temp1
    return new_theta
```

H Function

```
# Fonction hypothèse pour prédire les valeurs à partir des paramètres optimisés
def hyp(theta, indepX):
    return [theta[0] + theta[1] * x for x in indepX]
```

Defining Main Function

```
# Fonction principale pour entraîner le modèle et afficher les résultats
def main():
    # Initialisation des paramètres (theta) à zéro
    init_theta = np.zeros(2)
    learning_rate = 0.05 # Taux d'apprentissage pour contrôler la vitesse de convergence
    num_iterations = 56 # Nombre d'itérations pour optimiser les paramètres

    # Appel de la descente de gradient pour optimiser theta
    theta = gradientDescent(indepX_train, depY_train, init_theta, learning_rate, num_iterations)

    # Prédiction des valeurs avec les paramètres optimisés sur les données de test
    H = hyp(theta, indepX_test)

    #check
    for i in range(len(depY_test)):
        print(float(H[i]))
        #print(depY_test[i])
        print(depY_test.iloc[i])
        print('-----')
```


Main Function

```
# Point d'entrée du programme  
if __name__ == "__main__":  
    main()
```